



Content

- ☐ Data Preparation and Cleaning
- ☐ Train Validation Split 70-30
- ☐ EDA on Training Data
- ☐ Feature Engineering
- ☐ Model Building
- ☐ Predicting and Model Evaluation

Data Preparation and Cleaning

Import and inspect dataset

- Shape of dataset : (1000,40)

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 40 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   months_as_customer                       1000 non-null   int64
1   age                                       1000 non-null   int64
2   policy_number                           1000 non-null   int64
3   policy_bind_date                        1000 non-null   object
4   policy_state                            1000 non-null   object
5   policy_csl                              1000 non-null   object
6   policy_deductable                       1000 non-null   int64
7   policy_annual_premium                   1000 non-null   float64
8   umbrella_limit                          1000 non-null   int64
9   insured_zip                             1000 non-null   int64
10  insured_sex                             1000 non-null   object
11  insured_education_level                 1000 non-null   object
12  insured_occupation                     1000 non-null   object
13  insured_hobbies                         1000 non-null   object
14  insured_relationship                   1000 non-null   object
15  capital_gains                          1000 non-null   int64
16  capital_loss                           1000 non-null   int64
17  incident_date                          1000 non-null   object
18  incident_type                          1000 non-null   object
19  collision_type                          1000 non-null   object
20  incident_severity                      1000 non-null   object
21  authorities_contacted                   909 non-null    object
22  incident_state                         1000 non-null   object
23  incident_city                          1000 non-null   object
24  incident_location                      1000 non-null   object
25  incident_hour_of_the_day                1000 non-null   int64
26  number_of_vehicles_involved             1000 non-null   int64
27  property_damage                        1000 non-null   object
28  bodily_injuries                        1000 non-null   int64
29  witnesses                              1000 non-null   int64
30  police_report_available                 1000 non-null   object
31  total_claim_amount                     1000 non-null   int64
32  injury_claim                           1000 non-null   int64
33  property_claim                         1000 non-null   int64
34  vehicle_claim                          1000 non-null   int64
35  auto_make                              1000 non-null   object
36  auto_model                             1000 non-null   object
37  auto_year                              1000 non-null   int64
38  fraud_reported                         1000 non-null   object
39  c39                                     0 non-null      float64
dtypes: float64(2), int64(17), object(21)
memory usage: 312.6+ KB
```

Data Preparation and Cleaning

months_as_customer	0
age	0
policy_number	0
policy_bind_date	0
policy_state	0
policy_csl	0
policy_deductable	0
policy_annual_premium	0
umbrella_limit	0
insured_zip	0
insured_sex	0
insured_education_level	0
insured_occupation	0
insured_hobbies	0
insured_relationship	0
capital-gains	0
capital-loss	0
incident_date	0
incident_type	0
collision_type	0
incident_severity	0
authorities_contacted	91
incident_state	0
incident_city	0
incident_location	0
incident_hour_of_the_day	0
number_of_vehicles_involved	0
property_damage	0
bodily_injuries	0
witnesses	0
police_report_available	0
total_claim_amount	0
injury_claim	0
property_claim	0
vehicle_claim	0
auto_make	0
auto_model	0
auto_year	0
fraud_reported	0
c39	1000
dtype: int64	

Since nulls in `authorities_contacted` account for around 10% (91 rows), hence replaced them with 'Unknown' prevents potential data loss

authorities_contacted	
Police	292
Fire	223
Other	198
Ambulance	196
Unknown	91
Name: count, dtype: int64	

`_c39` contains only null values, so it was dropped from the dataset.

Redundant features like `policy_number`, `insured_zip`, `insured_hobbies`, and `incident_location` were dropped from the dataset.

Data Preparation and Cleaning

policy_bind_date	object
incident_date	object

Corrected the data type from object to datetime64[ns]

policy_bind_date	datetime64[ns]
incident_date	datetime64[ns]

Train Validation Split

```
# Put all the feature variables in X
X = df.drop('fraud_reported', axis=1)
# Put the target variable in y
y = df['fraud_reported']
```

Independent variable

Dependent variable





```
# Split the dataset into 70% train and 30% validation and use stratification on the target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=42)
# Reset index for all train and test sets
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
((700, 32), (700,), (300, 32), (300,))
```

Train size = 70%
Test size = 30%

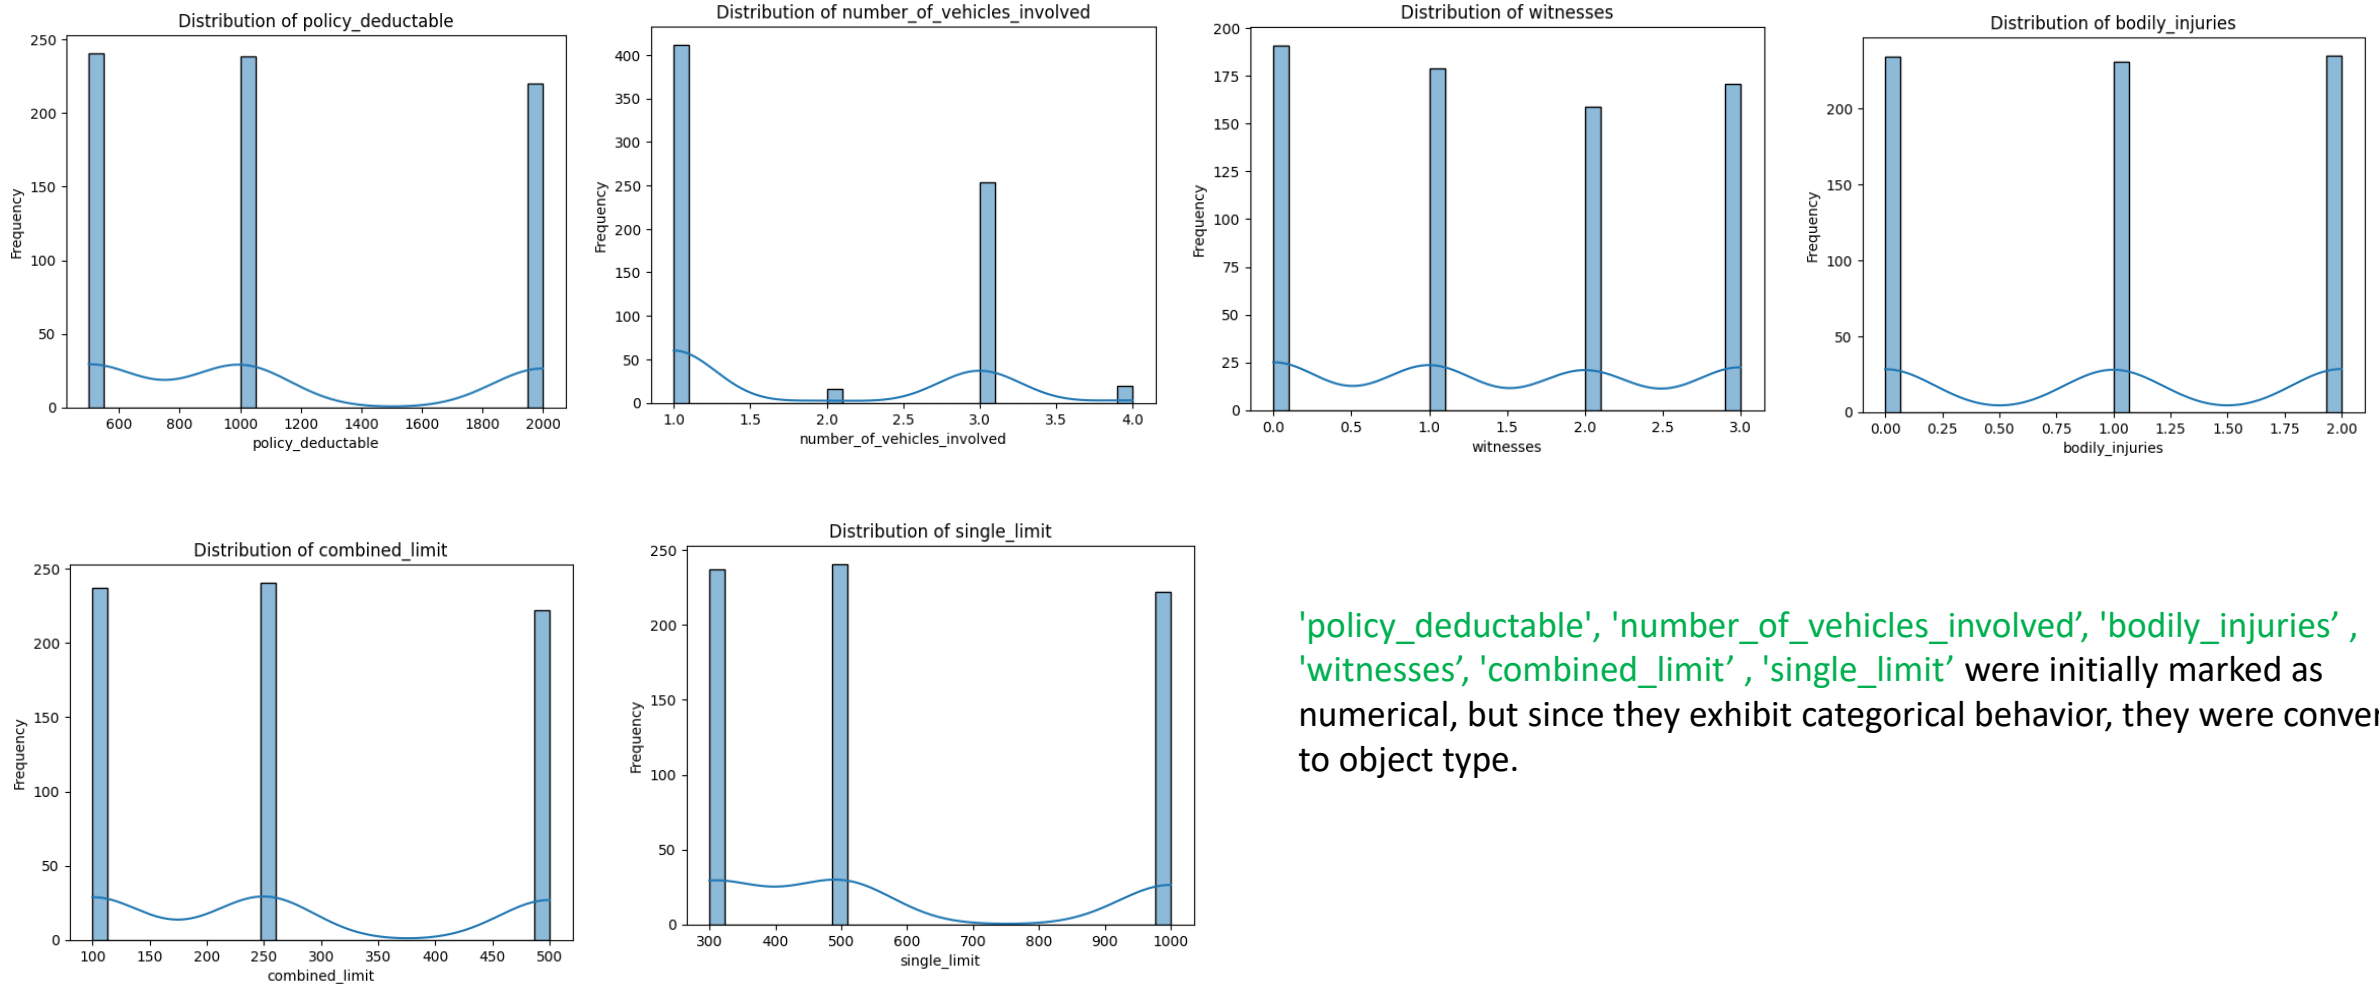
EDA on Numerical features

EDA on Training Data



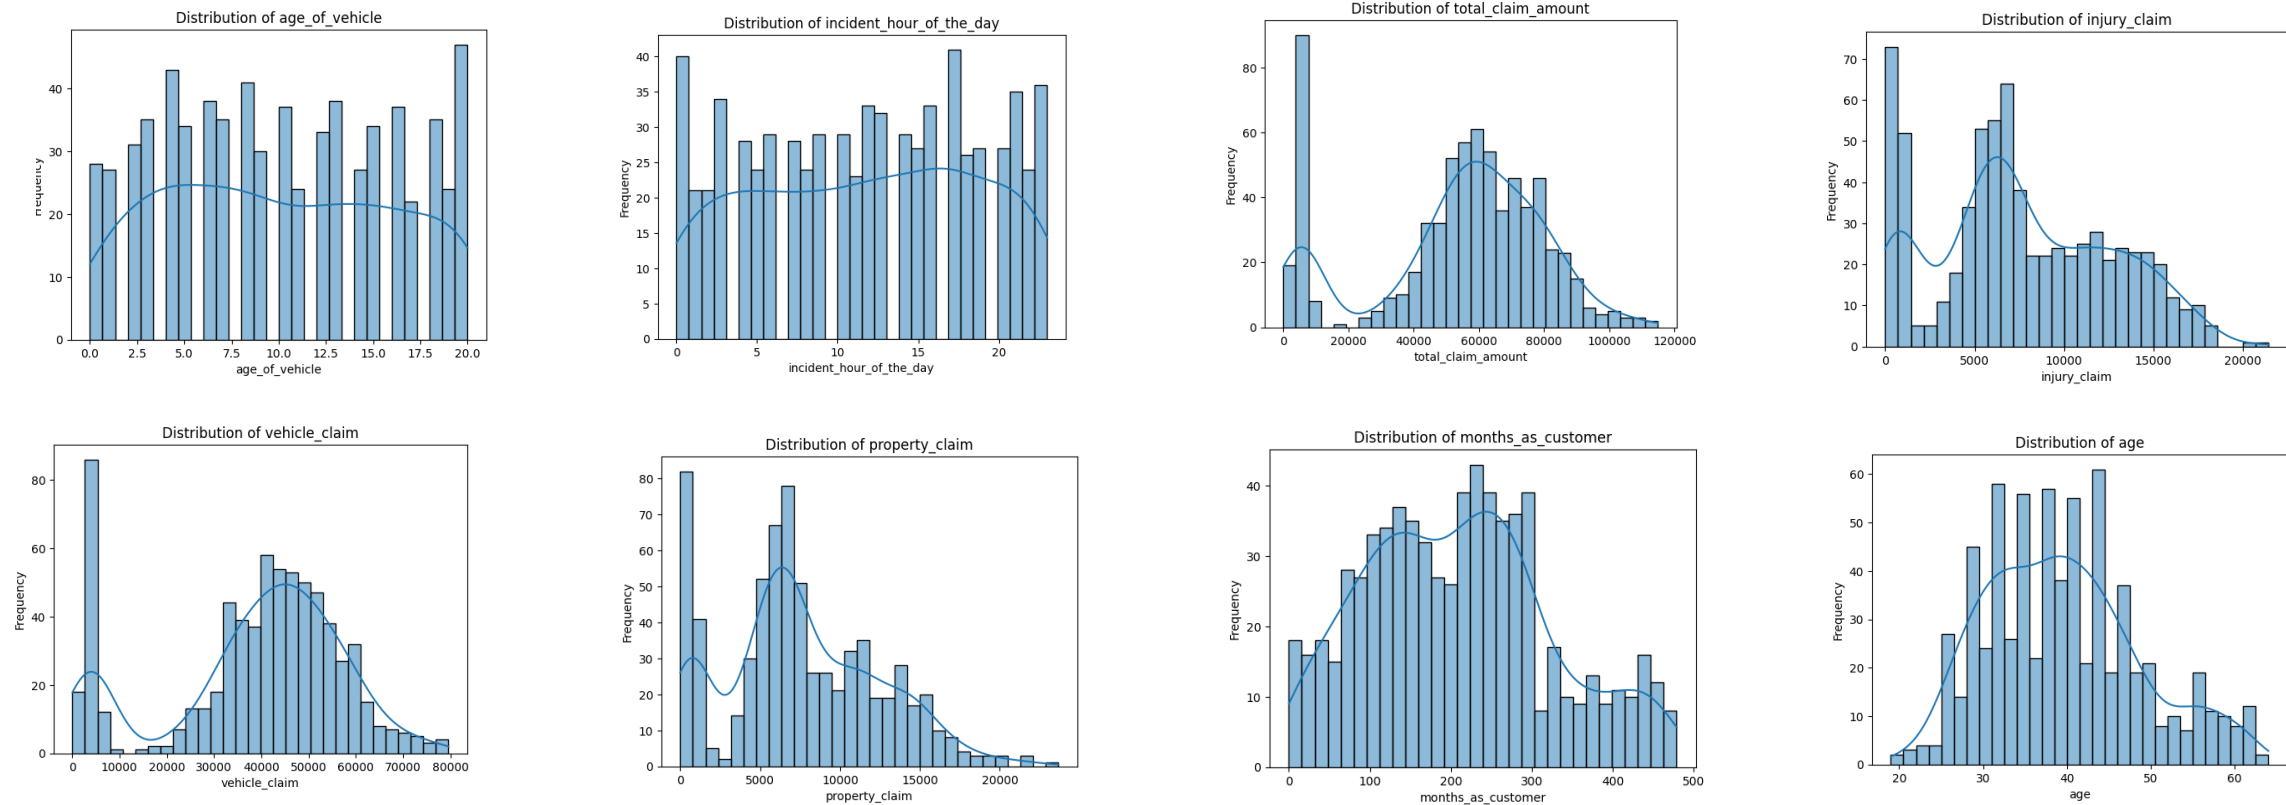
'umbrella_limit', 'capital-gains', 'capital-loss' are highly skewed toward zero, so they were deleted.

EDA on Training Data



'policy_deductable', 'number_of_vehicles_involved', 'bodily_injuries', 'witnesses', 'combined_limit', 'single_limit' were initially marked as numerical, but since they exhibit categorical behavior, they were converted to object type.

EDA on Training Data




'incident_hour_of_the_day', 'total_claim_amount', 'injury_claim', 'property_claim', 'age_of_vehicle', 'age' features are distributed across ranges, and to simplify analysis, these ranges have been grouped.

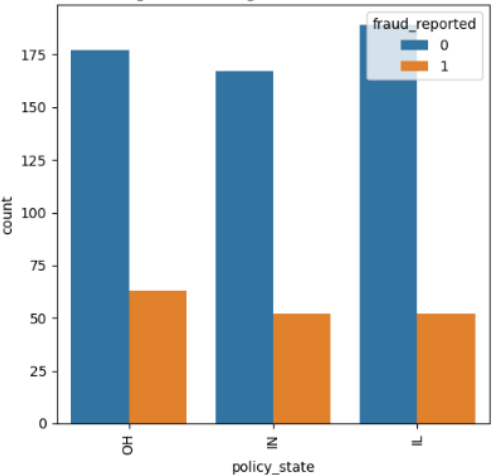
EDA on Categorical features

EDA on Training Data


Plotting data for the column: policy_state



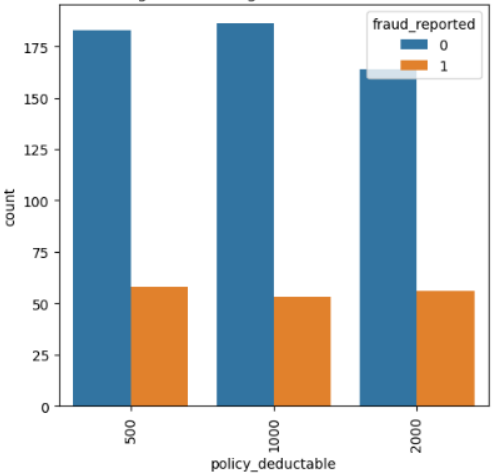
Plotting data for target in terms of total count




Plotting data for the column: policy_deductable



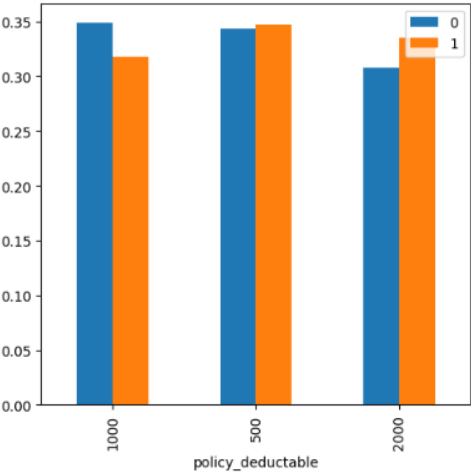
Plotting data for target in terms of total count



Plotting data for target in terms of percentage




Plotting data for target in terms of percentage



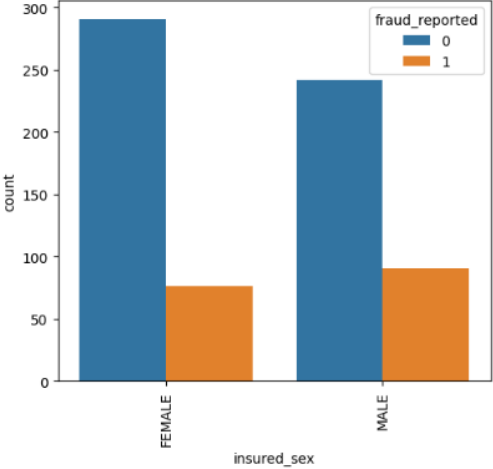
EDA on Training Data



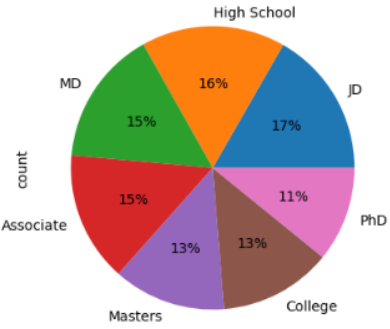
Plotting data for the column: insured_sex



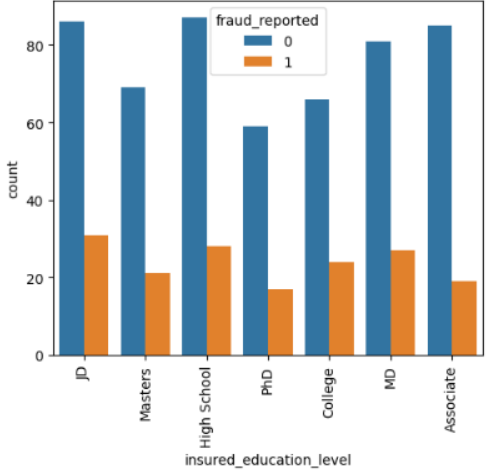
Plotting data for target in terms of total count



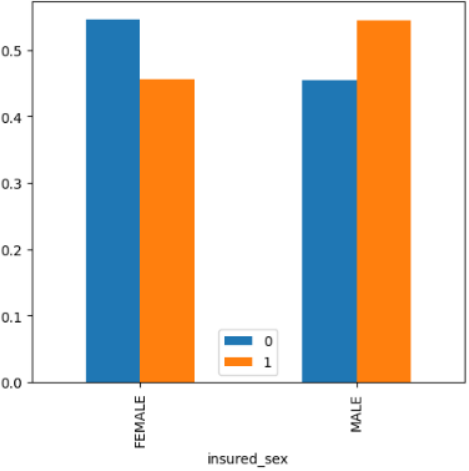
Plotting data for the column: insured_education_level




Plotting data for target in terms of total count



Plotting data for target in terms of percentage




Plotting data for target in terms of percentage

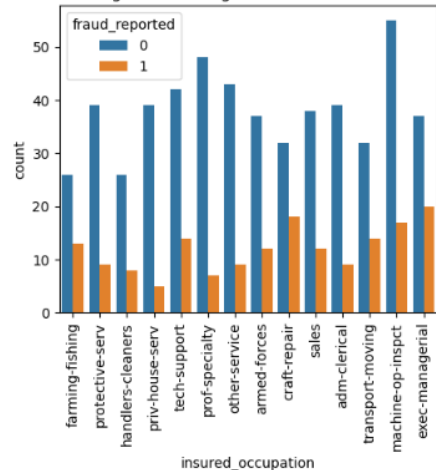


EDA on Training Data


Plotting data for the column: insured_occupation




Plotting data for target in terms of total count



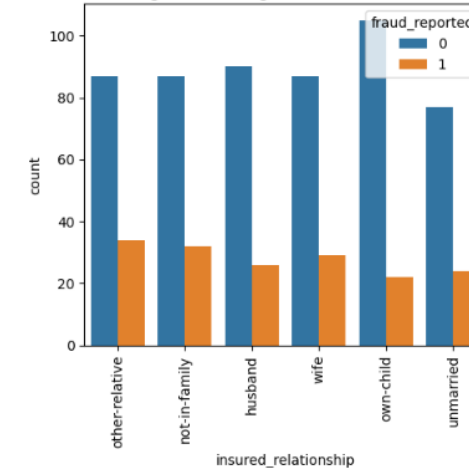
Plotting data for target in terms of percentage




Plotting data for the column: insured_relationship



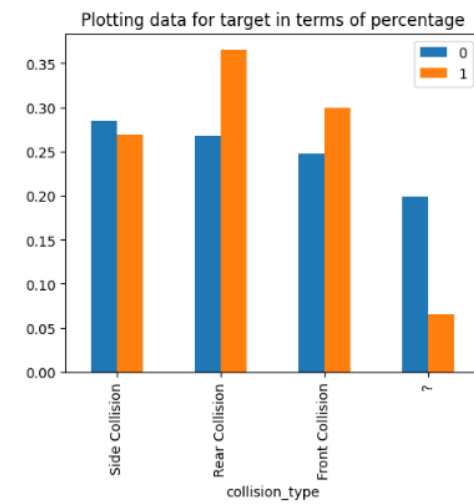

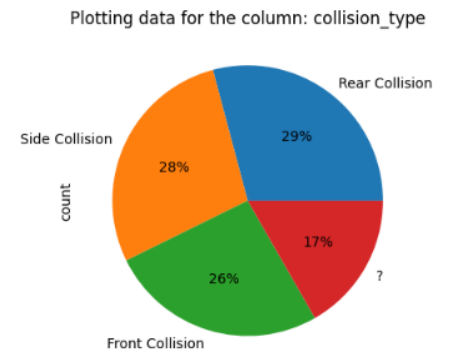
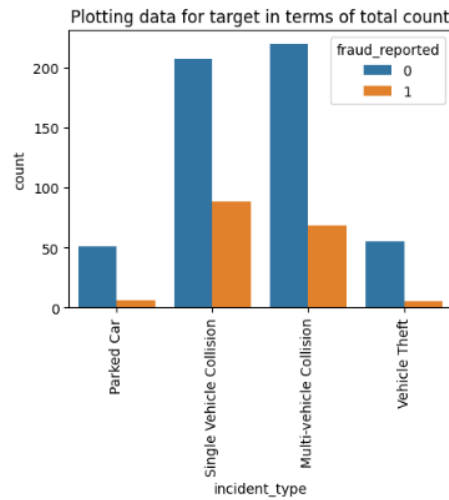

Plotting data for target in terms of total count





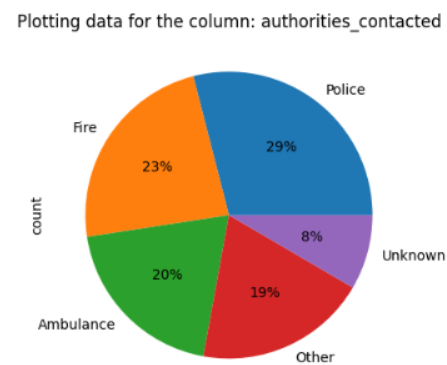



Plotting data for target in terms of percentage



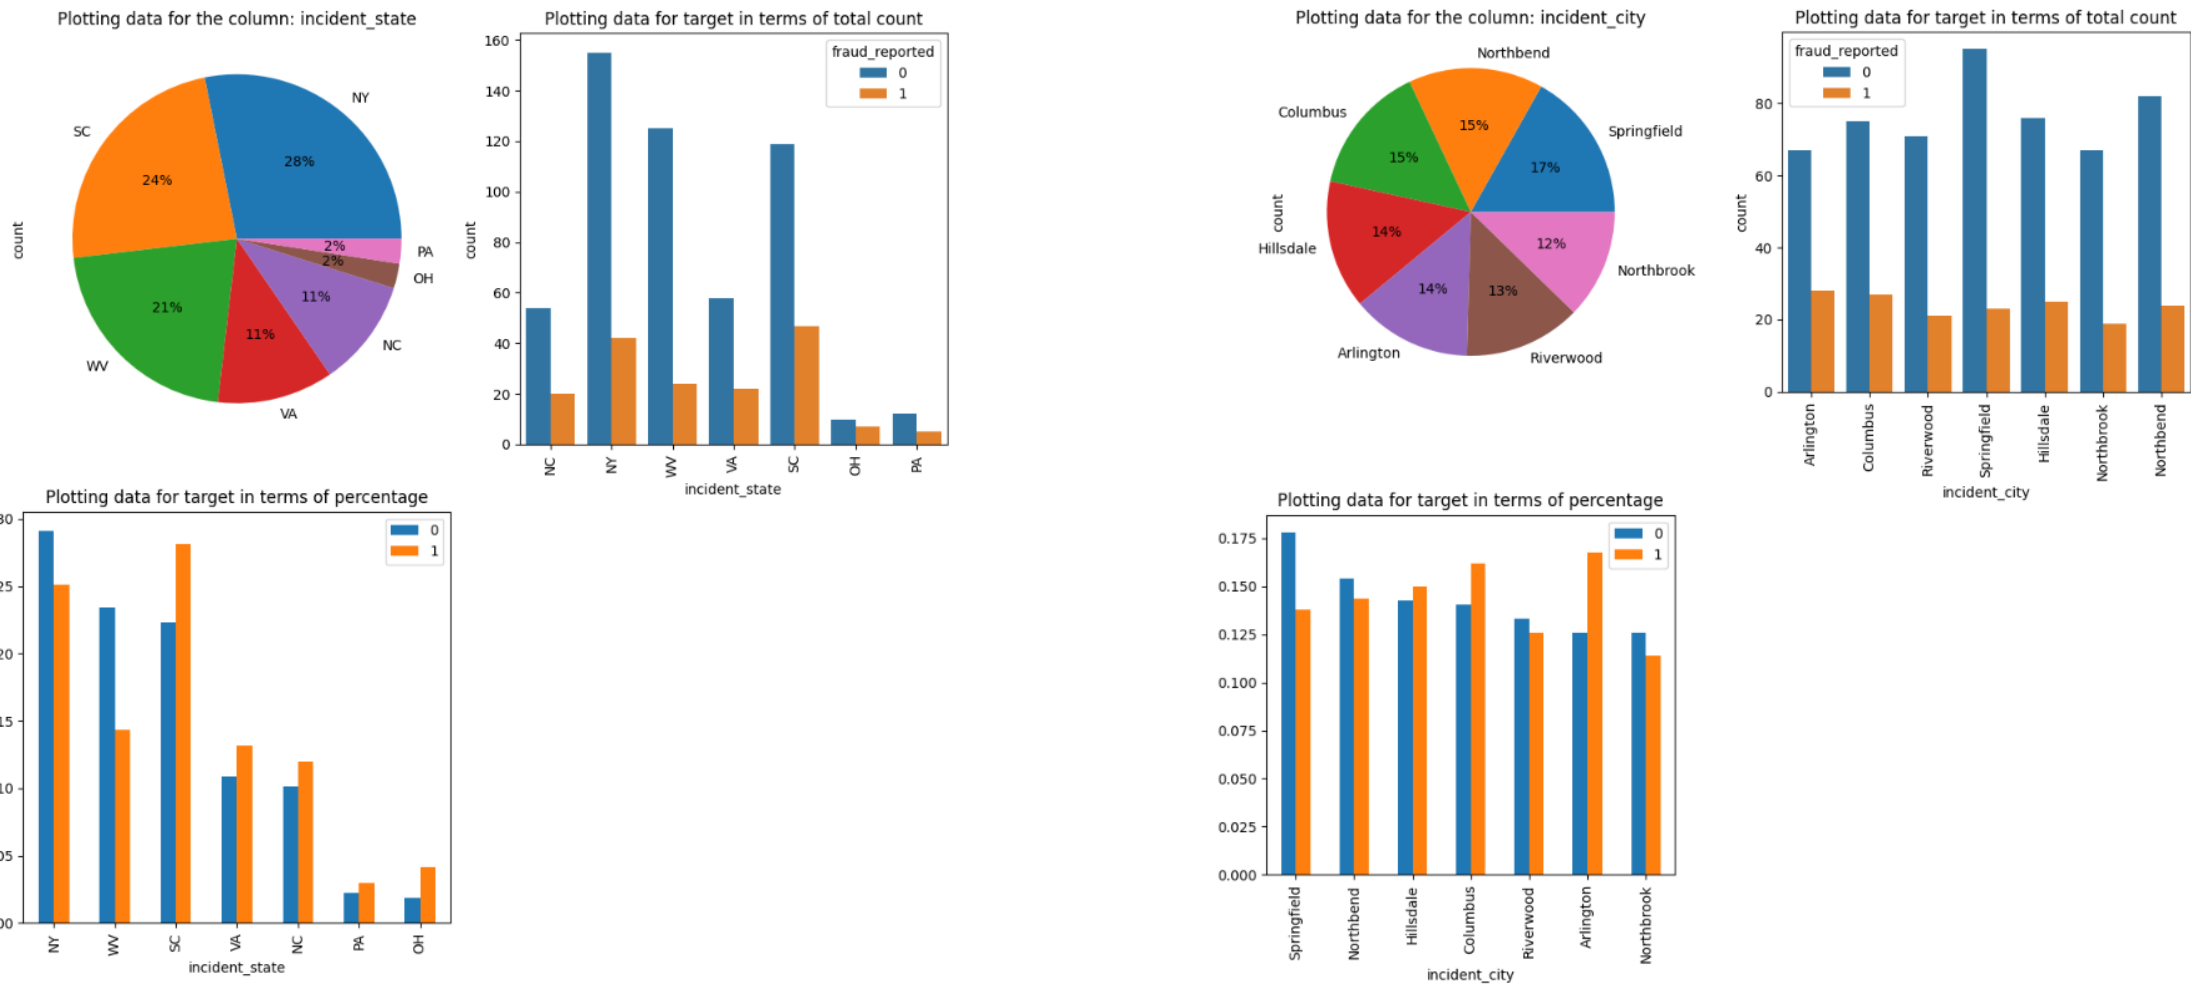
EDA on Training Data



EDA on Training Data




EDA on Training Data




EDA on Training Data




Plotting data for the column: number_of_vehicles_involved



Plotting data for target in terms of total count




Plotting data for target in terms of percentage




Plotting property_damage


Plotting data for the column: property_damage



Plotting data for target in terms of total count




Plotting data for target in terms of percentage




EDA on Training Data


Plotting data for the column: bodily_injuries




Plotting data for target in terms of total count




Plotting data for the column: witnesses




Plotting data for target in terms of total count



Plotting data for target in terms of percentage




Plotting data for target in terms of percentage




EDA on Training Data




Plotting data for the column: police_report_available




Plotting data for target in terms of total count




Plotting data for the column: auto_make




Plotting data for target in terms of total count




Plotting data for target in terms of percentage



Plotting data for target in terms of percentage




EDA on Training Data




EDA on Training Data




Plotting data for the column: time_category




Plotting data for target in terms of total count




Plotting data for the column: total_claim_range




Plotting data for target in terms of total count



Plotting data for target in terms of percentage




Plotting data for target in terms of percentage




EDA on Training Data




Plotting data for the column: injury_claim_range




Plotting data for target in terms of total count




Plotting data for the column: vehicle_claim_range




Plotting data for target in terms of total count



Plotting data for target in terms of percentage




Plotting data for target in terms of percentage




EDA on Training Data




Plotting data for the column: property_claim_range




Plotting data for target in terms of total count




Plotting data for the column: age_of_vehicle_range




Plotting data for target in terms of total count



Plotting data for target in terms of percentage




Plotting data for target in terms of percentage




EDA on Training Data




Plotting data for the column: age_of_client_range




Plotting data for target in terms of total count




Plotting data for the column: months_as_customer_range




Plotting data for target in terms of total count



Plotting data for target in terms of percentage



Plotting data for target in terms of percentage



Feature Engineering



policy_csl	single_limit	combined_limit
250/500	500	250
250/500	500	250
100/300	300	100
250/500	500	250
500/1000	1000	500

policy_csl contains combined limits. To facilitate analysis, we'll split it into separate columns and drop policy_csl feature


```
df['age_of_vehicle'] = df['incident_year'] - df['auto_year']
```

incident_year	auto_year	age_of_vehicle
2015	2004	11
2015	2007	8
2015	2007	8
2015	2014	1
2015	2009	6


Subtracting auto_year from incident_year provides the age_of_vehicle at the time of the incident

Model Building

Class balance check



Resampling



```
# Import RandomOverSampler from imblearn library
from imblearn.over_sampling import RandomOverSampler

# Perform resampling on training data
Over_sample = RandomOverSampler(random_state = 5)
X_resample_os, y_resample_os = Over_sample.fit_resample(X_train, y_train)
```

Model Building

Dummy variable creation and scaling

```
# Identify the categorical columns for creating dummy variables
categorical_cols = X_resample_os.select_dtypes(include=['object', 'category']).columns.tolist()
print('categorical_cols', categorical_cols)
```

```
# Create dummy variables using the 'get_dummies' for categorical columns in training data
dummy = pd.get_dummies(X_resample_os[categorical_cols], columns=categorical_cols, drop_first=True).astype('int')

X_resample_os = pd.concat([X_resample_os, dummy], axis=1)

X_resample_os.drop(categorical_cols, axis=1, inplace=True)
```

```
# Import the necessary scaling tool from scikit-learn
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

# Scale the numeric features present in the training data
X_resample_os[Numerical_cols] = scaler.fit_transform(X_resample_os[Numerical_cols])

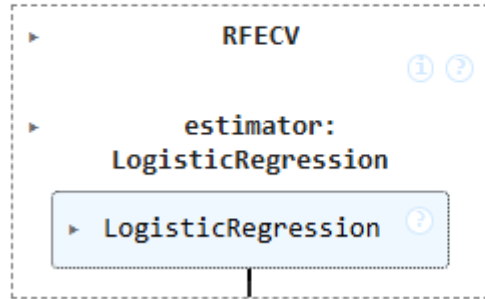
# Scale the numeric features present in the validation data
X_test[Numerical_cols] = scaler.fit_transform(X_test[Numerical_cols])
```

Model Building : Logistic regression

Feature selection

Apply RFECV to identify the most relevant features

```
rfecv = RFECV(estimator=LogisticRegression(solver='liblinear'), step=1, cv=StratifiedKFold(5), scoring='accuracy')  
rfecv.fit(X_train, y_train)
```



	Feature	Selected	Ranking
0	policy_annual_premium	True	1
1	policy_state_IN	True	1
2	policy_state_OH	True	1
4	policy_deductable_2000	True	1
6	insured_education_level_College	True	1
5	insured_sex_MALE	True	1
7	insured_education_level_High School	True	1
8	insured_education_level_JD	True	1
14	insured_occupation_exec-managerial	True	1
9	insured_education_level_MD	True	1
10	insured_education_level_Masters	True	1
12	insured_occupation_armed-forces	True	1
13	insured_occupation_craft-repair	True	1
19	insured_occupation_priv-house-serv	True	1
17	insured_occupation_machine-op-inspct	True	1
16	insured_occupation_handlers-cleaners	True	1
30	incident_type_Parked Car	True	1
25	insured_relationship_not-in-family	True	1
26	insured_relationship_other-relative	True	1
20	insured_occupation_prof-specialty	True	1

Model Building : Logistic regression

Best Logistic Regression Model

```
X_train_1 = X_train[top_features]
X_train_15 = sm.add_constant(X_train_1)
logm5 = sm.GLM(y_train,(sm.add_constant(X_train_1)),family = sm.families.Binomial())
logm5.fit().summary()
```

Generalized Linear Model Regression Results						
Dep. Variable:	fraud_reported	No. Observations:	1066			
Model:	GLM	Df Residuals:	1044			
Model Family:	Binomial	Df Model:	21			
Link Function:	Logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-688.48			
Date:	Sat, 14 Jun 2025	Deviance:	1377.0			
Time:	18:18:41	Pearson chi2:	1.08e+03			
No. Iterations:	4	Pseudo R-squ. (CS):	0.09026			
Covariance Type:	nonrobust					
	coef	std err	z	P> z	[0.025	0.975]
const	-0.0093	0.064	-0.145	0.885	-0.136	0.117
policy_annual_premium	-0.1441	0.066	-2.190	0.029	-0.273	-0.015
policy_state_IN	-0.0695	0.074	-0.939	0.348	-0.215	0.076
policy_state_OH	0.0622	0.075	0.830	0.407	-0.085	0.209
policy_deductable_2000	0.0609	0.066	0.924	0.355	-0.068	0.190
insured_education_level_College	0.1097	0.071	1.544	0.123	-0.030	0.249
insured_sex_MALE	0.2484	0.066	3.788	0.000	0.120	0.377
insured_education_level_High School	0.0665	0.073	0.913	0.361	-0.076	0.209
insured_education_level_JD	0.0833	0.073	1.144	0.253	-0.059	0.226
insured_occupation_exec-managerial	0.1091	0.069	1.573	0.116	-0.027	0.245
insured_education_level_MD	0.0501	0.072	0.693	0.488	-0.092	0.192
insured_occupation_armed-forces	0.0490	0.066	0.741	0.459	-0.081	0.179
insured_occupation_craft-repair	0.1532	0.069	2.227	0.026	0.018	0.288
insured_occupation_priv-house-serv	-0.2364	0.076	-3.129	0.002	-0.384	-0.088
incident_type_Parked Car	-0.1985	0.074	-2.686	0.007	-0.343	-0.054
insured_relationship_not-in-family	0.1349	0.068	1.988	0.047	0.002	0.268
insured_relationship_other-relative	0.2079	0.068	3.057	0.002	0.075	0.341
insured_occupation_prof-specialty	-0.1385	0.069	-2.005	0.045	-0.274	-0.003
insured_occupation_sales	0.0473	0.067	0.702	0.482	-0.085	0.179
insured_occupation_protective-serv	-0.0560	0.067	-0.837	0.402	-0.187	0.075
insured_occupation_transport-moving	0.1062	0.067	1.580	0.114	-0.026	0.238
incident_type_Single Vehicle Collision	0.2418	0.067	3.583	0.000	0.110	0.374

	Feature	VIF
2	policy_state_OH	1.36
1	policy_state_IN	1.33
6	insured_education_level_High School	1.29
9	insured_education_level_MD	1.27
7	insured_education_level_JD	1.26
4	insured_education_level_College	1.23
8	insured_occupation_exec-managerial	1.16
19	insured_occupation_transport-moving	1.12
17	insured_occupation_sales	1.12
11	insured_occupation_craft-repair	1.12
16	insured_occupation_prof-specialty	1.12
15	insured_relationship_other-relative	1.12
14	insured_relationship_not-in-family	1.12
20	incident_type_Single Vehicle Collision	1.11
10	insured_occupation_armed-forces	1.10
18	insured_occupation_protective-serv	1.10
13	incident_type_Parked Car	1.09
12	insured_occupation_priv-house-serv	1.07
0	policy_annual_premium	1.05
3	policy_deductable_2000	1.05
5	insured_sex_MALE	1.04


Model Building : Logistic regression



```
y_train_pred_final['Predicted'] = y_train_pred_final.Predicted_Prob.map(lambda x: 1 if x > 0.5 else 0)
```

[322, 211]
[185, 348]

Metrics	value	Remarks
accuracy	0.628517824	62.85%
sensitivity/Recall	0.652908068	65.3% of actual positives were correctly predicted
specificity	0.60412758	60.4% of actual negatives were correctly predicted
Precision	0.62254025	62.2% of predicted positives were correct
f1_score	0.637362637	Balance between precision and recall




Model Building : Logistic regression

```
numbers = [float(x)/10 for x in range(10)]  
for i in numbers:  
    y_train_pred_final[i]= y_train_pred_final.Predicted_Prob.map(lambda x: 1 if x > i else 0)
```

	Actual	Predicted_Prob	Predicted	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9
0	0	0.313163	0	1	1	1	1	0	0	0	0	0	0
1	0	0.578618	1	1	1	1	1	1	1	0	0	0	0
2	1	0.569165	1	1	1	1	1	1	1	0	0	0	0
3	0	0.582130	1	1	1	1	1	1	1	0	0	0	0
4	0	0.294940	0	1	1	1	0	0	0	0	0	0	0

	prob	accuracy	sensi	speci
0.0	0.0	0.500000	1.000000	0.000000
0.1	0.1	0.500938	0.996248	0.005629
0.2	0.2	0.517824	0.986867	0.048780
0.3	0.3	0.556285	0.939962	0.172608
0.4	0.4	0.593809	0.848030	0.339587
0.5	0.5	0.628518	0.652908	0.604128
0.6	0.6	0.590994	0.354597	0.827392
0.7	0.7	0.542214	0.114447	0.969981
0.8	0.8	0.504690	0.011257	0.998124
0.9	0.9	0.500000	0.000000	1.000000




Model Building : Logistic regression



```
# Create a column for final prediction based on the optimal cutoff
y_train_pred_final['Optimal_Predicted'] = y_train_pred_final.Predicted_Prob.map(lambda x: 1 if x > 0.45 else 0)
```

Metrics	value	Remarks
accuracy	0.614446529	61.44%
sensitivity/Recall	0.772983114	77.3% of actual positives were correctly predicted
specificity	0.455909944	45.6% of actual negatives were correctly predicted
Precision	0.586894587	58.7% of predicted positives were correct
f1_score	0.667206478	Balance between precision and recall

precision-recall curve



Model Building : Random forest

```
# Build a base random forest model
rf = RandomForestClassifier(n_estimators=100, max_depth=4, max_features=5, random_state=100, oob_score=True)
rf.fit(X_train, y_train)
```

RandomForestClassifier

RandomForestClassifier(max_depth=4, max_features=5, oob_score=True, random_state=100)

	features	Imp
36	incident_severity_Minor Damage	0.088142
37	incident_severity_Total Loss	0.057772
42	authorities_contacted_Unknown	0.036399
38	incident_severity_Trivial Damage	0.035396
0	policy_annual_premium	0.033319
35	collision_type_Unknown	0.024426
97	vehicle_claim_range_40k-60k	0.023208
27	insured_relationship_own-child	0.021713
26	insured_relationship_other-relative	0.020784
79	combined_limit_500	0.019508
5	insured_sex_MALE	0.019384
30	incident_type_Parked Car	0.019108
32	incident_type_Vehicle Theft	0.017087
19	insured_occupation_priv-house-serv	0.015043
100	property_claim_range_5k-10k	0.014860

Metrics	value	Remarks
accuracy	0.815196998	81.51%
sensitivity/Recall	0.791744841	79.17% of actual positives were correctly predicted
specificity	0.838649156	83.9% of actual negatives were correctly predicted
Precision	0.830708661	83.1% of predicted positives were correct
f1_score	0.810758886	Balance between precision and recall

```
# Use cross validation to check if the model is overfitting
cv_scores = cross_val_score(rf, X_train_rf, y_train, cv=5, scoring='accuracy')

print("Cross-Validation Accuracy Scores:", cv_scores)
print("Mean CV Accuracy:", cv_scores.mean())
```

Cross-Validation Accuracy Scores: [0.72897196 0.7370892 0.79342723 0.82629108 0.84507042]
Mean CV Accuracy: 0.7861699793778246

Training Accuracy 81.51%

Mean Cross-Validation Accuracy 78.61%

The accuracy gap is just 2.9% — which is small and acceptable and model is NOT significantly overfitting

Model Building : Random forest

Hyperparameter Tuning

```
# Use grid search to find the best hyperparamter values
Classifier_rf = RandomForestClassifier(random_state=42, n_jobs=-1)
# Best Hyperparameters
params = {
    'max_depth': [1, 2, 5, 10, 20],
    'min_samples_leaf': [5, 10, 20, 50, 100],
    'max_features': [2,3,4],
    'n_estimators': [10, 30, 50, 100, 200]
}
```

```
# Building random forest model based on results of hyperparameter tuning
grid_search = GridSearchCV(estimator=Classifier_rf, param_grid= params,cv=4, n_jobs=-1, verbose=1, scoring='accuracy')
```

```
# Make predictions on training data
grid_search.fit(X_train_rf,y_train)
```

Fitting 4 folds for each of 375 candidates, totalling 1500 fits

GridSearchCV

best_estimator_: RandomForestClassifier

RandomForestClassifier

RandomForestClassifier(max_depth=20, max_features=4, min_samples_leaf=5, n_estimators=30, n_jobs=-1, random_state=42)

Metrics	value	Remarks
accuracy	0.873358349	87.30%
sensitivity/Recall	0.853658537	85.36% of actual positives were correctly predicted
specificity	0.889305816	88.9% of actual negatives were correctly predicted
Precision	0.885214008	88.5% of predicted positives were correct
f1_score	0.869149952	Balance between precision and recall

Model Evaluation



Logistic regression

Metrics	value	Remarks
accuracy	0.403333333	40.30%
sensitivity/Recall	0.75	75 % of actual positives were correctly predicted
specificity	0.277272727	27.7% of actual negatives were correctly predicted
Precision	0.273972603	27.4% of predicted positives were correct
f1_score	0.401337793	Balance between precision and recall

Random forest

Metrics	value	Remarks
accuracy	0.773333333	77.30%
sensitivity/Recall	0.75	75% of actual positives were correctly predicted
specificity	0.8227	82.3% of actual negatives were correctly predicted
Precision	0.5666	56.6% of predicted positives were correct
f1_score	0.6	Balance between precision and recall

Conclusion

Logistic regression

Metrics	Train Data	Test Data
Accuracy	0.6144	0.4033
Sensitivity/Recall	0.7730	0.7500
Specificity	0.4559	0.2773
Precision	0.5869	0.2740
f1_score	0.6672	0.4013

- The logistic regression model is overfitting — performing reasonably well on the training data but very poorly on the test set.
- The recall remains high, which means the model still finds most positive cases, but it sacrifices precision and specificity, resulting in many false positives.
- Model is unreliable for deployment in its current form

Random forest:

Metrics	Train Data	Test Data
Accuracy	0.8734	0.7733
Sensitivity/Recall	0.8537	0.7500
Specificity	0.8893	0.8227
Precision	0.8852	0.5666
f1_score	0.8691	0.6000

- The Random Forest model demonstrates strong and well-rounded performance on both training and test datasets.
- It generalizes fairly well with a good balance of sensitivity (recall) and specificity, making it reliable for binary classification tasks.
- While test precision is lower, it's still usable — but we may need to adjust thresholds depending on business needs (e.g., favoring precision over recall or vice versa).
- Overall, this model is a good candidate for deployment or further tuning.